

EMOVO Speech Emotion Recognition Project

Nikolaos Mouzakitis MTP321

April 5, 2025

1 Introduction

This document details the implementation of a speech emotion recognition system using the EMOVO dataset. The system employs visibility graph-based features and machine learning classifiers for gender and emotion recognition.

2 System Architecture

The system consists of three main components:

- Feature extraction pipeline (source: `part1_extract_features.py`)
- Gender classification module (source: `part2_svm_rf_gender_classification.py`)
- Emotion classification module (source: `part3_emotion_classification.py`)

3 Feature Extraction Process

3.1 Methodology

In the feature extraction process, the database's audio signals are getting converted into graph-based features through several steps:

3.1.1 Audio Loading and Preprocessing

Initially the dataset, is read and by using the naming format speaker and emotion is registered for each of the samples.

```
1 def load_emovo_dataset(root_dir):
2     emotion_map = {
3         'dis': 'disgust', 'gio': 'joy', 'pau': 'fear',
4         'rab': 'anger', 'sor': 'surprise', 'tri': 'sadness', 'neu':
5         'neutral'
6     }
7     data = []
8     for speaker in os.listdir(root_dir):
```

```

9     speaker_dir = os.path.join(root_dir, speaker)
10    if os.path.isdir(speaker_dir):
11        for file in os.listdir(speaker_dir):
12            if file.endswith(".wav"):
13                parts = file.split("-")
14                if len(parts) >= 3:
15                    emotion_code = parts[0]
16                    data.append({
17                        "speaker": speaker,
18                        "emotion": emotion_map.get(emotion_code
19, "unknown"),
20                        "filepath": os.path.join(speaker_dir,
21file)
22                    })
23    return pd.DataFrame(data)

```

Key aspects:

- Mappings of EMOVO’s emotion codes to full names
- Organization of the files by speaker and emotion
- Handles directory traversal safely by ensuring we read only the `wav` files.

3.1.2 Sliding Window Processing

In the next step, three different sliding window lengths were examined (1000 (*overlap:500*), 2000 (*overlap:1000*), 3000 (*overlap:1000*)).

```

1 def sliding_window_std(audio, window_size=1500, overlap=900):
2     std_values = []
3     hop = window_size - overlap
4     for i in range(0, len(audio) - window_size + 1, hop):
5         std_values.append(np.std(audio[i:i + window_size]))
6     return np.array(std_values)

```

Parameters:

- `window_size`: Number of samples in each window (tested 1000, 2000, 3000)
- `overlap`: Samples shared between consecutive windows (500 samples for the 1000 sliding window and overlap 1000 for the rest)
- Standard deviation within each window is computed.

3.2 Visibility Graph Construction in EMOVO Processing

In this implementation, the visibility graph is constructed from the standard deviation array(not the zero-mean RMS energy as in the related given paper [1]), which is generated through sliding window analysis of audio signals. The complete transformation pipeline works as follows:

3.2.1 Input Preparation

```
1 # From part1_extract_features.py
2 T = sliding_window_std(audio, window_size=2000, overlap=1000)
```

The process begins with:

- Audio signal divided into samples of the length of the sliding window
- overlap (500 or 1000 samples in our case) between consecutive windows
- Standard deviation computed for each window, creating array $T = [\sigma_1, \sigma_2, \dots, \sigma_n]$

3.2.2 Graph Construction

```
1 def compute_visibility_graph(T):
2     vg = ts2vg.NaturalVG()
3     vg.build(T)
4     return vg.adjacency_matrix()
```

For our speech emotion analysis:

- Each σ_i becomes node v_i in the graph
- Edges connect nodes where the σ values are mutually visible
- Visibility condition: $\sigma_k < \sigma_j + (\sigma_i - \sigma_j) \frac{j-k}{j-i}$ for all k between i and j

3.2.3 Example of visibility graph adjacency matrix

Consider a test sample from a speaker's recording std devs:

$$T = [0.21, 0.35, 0.18, 0.42, 0.29] \quad (1)$$

The visibility graph adjacency matrix becomes:

$$A = \begin{bmatrix} 0 & 1 & 0 & 1 & 0 \\ 1 & 0 & 1 & 1 & 1 \\ 0 & 1 & 0 & 1 & 0 \\ 1 & 1 & 1 & 0 & 1 \\ 0 & 1 & 0 & 1 & 0 \end{bmatrix} \quad (2)$$

In a visibility graph a time series or data generated from time series are represented as a graph, with nodes corresponding on these time points while the edges are connected only if nodes "see" each other.

3.2.4 Graph Feature Extraction

The adjacency matrix A generated by the visibility graph undergoes several transformations to extract emotion-relevant features. For each audio segment, we process A as follows:

3.2.5 Matrix Preprocessing

```

1 if isinstance(A, np.ndarray):
2     G = nx.from_numpy_array(A)
3     avg_A = np.mean(A)
4 else:
5     G = nx.from_scipy_sparse_array(A)
6     avg_A = A.mean()

```

- **Matrix Conversion:** Handles both dense and sparse matrix representations
- **NetworkX Graph:** Converts A into a graph object G for topological analysis
- **Mean Calculation:** Computes $\text{Avg_A} = \frac{1}{n^2} \sum_{i,j} A_{ij}$ as baseline connectivity measure

3.2.6 Feature Computation Pipeline

For each audio sample, six key features are extracted as instructed by the project rules:

1. Degree of Connectivity (DoC):

$$\text{DoC} = \sum_{i=1}^n \sum_{j=1}^n A_{ij} \quad (3)$$

- Measures total graph connectivity

2. Graph Density:

$$\text{Density} = \frac{\text{DoC}}{n(n-1)} \quad (4)$$

- Ratio of existing edges to possible edges

3. Clustering Coefficient (CC):

$$\text{CC} = \frac{1}{n} \sum_{i=1}^n \frac{2T_i}{k_i(k_i-1)} \quad (5)$$

- T_i : Triangles around node i , k_i : degree of node i
- Measures local connectivity

4. Average Path Length (L):

$$L = \frac{1}{n(n-1)} \sum_{i \neq j} d(v_i, v_j) \quad (6)$$

- $d(v_i, v_j)$: Shortest path between nodes

- Only computed for connected components

5. Modularity (M):

$$M = \frac{1}{2m} \sum_{ij} \left[A_{ij} - \frac{k_i k_j}{2m} \right] \delta(c_i, c_j) \quad (7)$$

- m : Total edges, k_i : Node degree, c_i : Community

6. Average Adjacency (Avg_A):

$$\text{Avg_A} = \text{mean}(A) \quad (8)$$

- (Extra feature that project mentions to add) Baseline connectivity measure, serves as a normalized version of DoC

```

1 def extract_graph_features(A):
2     if isinstance(A, np.ndarray):
3         G = nx.from_numpy_array(A)
4         avg_A = np.mean(A)
5     else:
6         G = nx.from_scipy_sparse_array(A)
7         avg_A = A.mean()
8
9     features = {
10         "DoC": np.sum(A) if isinstance(A, np.ndarray) else A.sum(),
11         "Density": nx.density(G),
12         "CC": nx.average_clustering(G),
13         "L": nx.average_shortest_path_length(G) if nx.is_connected(G) else np.nan,
14         "M": compute_modularity(G),
15         "Avg_A": avg_A
16     }
17     return features

```

4 Gender Classification

4.1 Implementation Details

```

1 def evaluate_model(model, X, y, groups):
2     logo = LeaveOneGroupOut()
3     male_true, male_pred = [], []
4     female_true, female_pred = [], []
5
6     for fold, (train_idx, test_idx) in enumerate(logo.split(X, y, groups)):
7         X_train, X_test = X[train_idx], X[test_idx]
8         y_train, y_test = y[train_idx], y[test_idx]
9
10        model.fit(X_train, y_train)
11        y_pred = model.predict(X_test)
12

```

```

13     if y_test[0] == 1: # Male
14         male_true.extend(y_test)
15         male_pred.extend(y_pred)
16     else: # Female
17         female_true.extend(y_test)
18         female_pred.extend(y_pred)

```

In the gender classification question, the method of Leave-One-Speaker-Out cross-validation was used as instructed with separate tracking of male/female performance.

4.2 Classification Results




Figure 1: Comparisson of features sliding window 1000




Figure 2: Comparisson of features sliding window 2000




Figure 3: Comparisson of features sliding window 3000

As we can observe in the comparative distribution plots above, the six features obtained are not enough and do not possess the discriminative ability to separate in an acceptable manner gender in our case.

The average results of all the three selected combinations are presented in the following table. Best classification results were obtained by the SVM method for the window size of 2000.

Gender	Classifier	Accuracy	Sensitivity	Specifity	F1-Score
Male	SVM	25.964%	0.26	0.453	0.412
Female	SVM	45.238%	0.453	0.26	0.612
Male	RF	21.33%	0.213	0.328	0.612
Female	RF	32.766%	0.328	0.213	0.493

Table 1: Gender classification average performance

5 Emotion Classification

5.1 Implementation Details

```

1 # Random Forest with GridSearch
2 param_grid_rf = {
3     'n_estimators': [50, 100, 200],
4     'max_depth': [None, 10, 20],
5     'class_weight': ['balanced', None],
6     'max_features': ['sqrt', 'log2']
7 }
8
9 # SVM with RandomizedSearch
10 param_dist_svm = {
11     'kernel': ['linear', 'rbf'],
12     'C': [0.1, 1, 10],
13     'gamma': ['scale', 0.01],
14     'class_weight': [None, 'balanced']
15 }
16
17 for train_idx, test_idx in logo.split(X, y, groups):
18     X_train, X_test = X.iloc[train_idx], X.iloc[test_idx]
19     y_train, y_test = y.iloc[train_idx], y.iloc[test_idx]
20
21     # Apply SMOTE only to training data
22     X_train_res, y_train_res = smote.fit_resample(X_train, y_train)
23
24     # Scale features for SVM
25     scaler = StandardScaler()
26     X_train_scaled = scaler.fit_transform(X_train_res)
27     X_test_scaled = scaler.transform(X_test)

```

For the emotion classification task, hyperparameter optimization for both classifiers was conducted in order to improve the results (f.e. utilization of SMOTE for class imbalance (applied only to training data). Feature scaling was used in the SVM case.

5.2 Emotion Classification Results

The average results of the classification of emotions (source: `part3_emotion_classification.py`), after gathering outputs of the program's execution for the three selected combinations of sliding windows and overlaps are presented in the following table.

Table 2: Per-Emotion Classification Performance (Averaged across speakers)

Emotion	SVM				RandomForest			
	Acc.	Sens.	Spec.	F1	Acc.	Sens.	Spec.	F1
Anger	0.746	0.197	0.837	0.186	0.743	0.221	0.829	0.197
Fear	0.807	0.083	0.928	0.098	0.767	0.102	0.878	0.111
Disgust	0.754	0.23	0.841	0.182	0.753	0.221	0.842	0.198
Joy	0.783	0.199	0.897	0.106	0.764	0.122	0.871	0.13
Neutral	0.707	0.372	0.762	0.253	0.779	0.277	0.863	0.249
Sadness	0.813	0.154	0.922	0.162	0.785	0.126	0.894	0.139
Surprise	0.754	0.154	0.854	0.118	0.769	0.201	0.864	0.190
Avg	0.766	0.198	0.863	0.158	0.766	0.182	0.863	0.173

Acc. = Accuracy, **Sens.** = Sensitivity, **Spec.** = Specificity

6 Conclusion

The implemented system demonstrates that graph-based features extracted from speech signals can effectively capture both gender and emotion information. Key findings include:

- Optimal window size of 2000 samples
- SVM outperforms Random Forest for both tasks
- Graph features particularly effective for neutral and anger emotions
- Speaker-independent evaluation shows robustness

6.1 Other notes

Graph-based features could be used in conjunction alongside other acoustic features from the samples (f.e *pitch_std*) in order to obtain better gender classification results, since acoustic features appear to offer better discrimination ability for gender selection.

7 References

- A. Pentari, G. Kafentzis, and M. Tsiknakis, "Investigating graph-based features for speech emotion recognition," tech. rep., FORTH-ICS and University of Crete, 2023.